Second Look — live demo sheet (local app, AI on)

About 2 minutes, everything on **http://localhost:5173** with the four agents on qwen3-8b via Foundry Local. Video: under 2:00, YouTube **Public**. Deadline: 27 Oct, 11:59 pm Pacific (2:59 am Eastern).

The trick

• **Start the AI question first** (Investigate this decision on The press), then spend its 1–3 minutes on the instant pages. The question keeps running while you change pages.

• Never refresh while an answer is running. One question at a time.

• If the answer isn't ready when you return: play the move on The press, or open Also in this passage and click a stat. In a video, cut the wait and label it "≈1 min later — sped up".

Setup (10 minutes before)

• `foundry server start`, then `foundry model load qwen3-8b`.

• Backend (in backend): `.\.venv\Scripts\python.exe -m uvicorn app.main:app --host 127.0.0.1 --port 8000`. Frontend (in frontend): `npm run dev`.

• Header must read *Foundry Local (on-device) · qwen3-8b · Agent Framework*.

• Warm up: on The argument click Pressing decision once and let it finish; refresh; open The press.

• Browser about 1600×1000, zoom 100%, bookmarks hidden, notifications off, GPU-heavy apps closed.

Don't

• Call the on-device model "cloud Foundry", or say the agents write the explanation (they choose and order the evidence; code writes the facts).

• Claim personalization, real-time analysis, multiple languages, real players or real-match accuracy. No club logos, real footage or copyrighted music.

The script

Time	Page	On screen	Say (suggested)
0:00–0:20	The press	Question, "No. Holding was safer.", 30 vs 71, orange 8.5 m gap. Click **Investigate this decision** (shows "planning…").	"Should the right-back have pressed earlier? Second Look replays the next eight seconds two hundred times for each choice. If he holds, Harbor reach a dangerous position in 30 of 200. If he presses earlier, 71. Lower is safer: holding wins. I've just handed the same question to four AI agents running on this laptop; they'll take a minute or two."
0:20–0:50	Three choices	Hold is marked. Drag **Cover delay** 0.6 s → 1.2 s; headline becomes "Holding position is the clear call".	"With a teammate covering the inside pass it's too close to call, under a rule we fixed before any run. Make that cover arrive 0.6 seconds later and holding becomes the clear call. The answer depends on that assumption, so the app shows it."
0:50–1:05	Three choices	**Run stress test**; show the sentence. Optional: Watch an example.	"A challenger searches a grid of assumptions for settings that flip the answer. The agents run this same search."
1:05–1:35	The argument	Answered: badge **AI agents · on-device · qwen3-8b**; expand **How this answer was made**; scroll to a validator line with rejected claims.	"Four Microsoft Agent Framework agents — a planner, an evidence analyst, a challenger and a producer — running on this laptop with Foundry Local. The agents choose the evidence and order it, but code writes every number. When the producer wrote a claim the numbers don't support, the validator rejected it."
1:35–1:45	The argument	Switch to **Fan view** (instant).	"Same facts for a fan: plain words, the caveat attached, nothing re-run."
1:45–2:00	On air	**Play recap**; end on the SIMULATION card.	"For a broadcast desk, the agents' answer becomes timed cards on the replay. Every card cites a fact, and simulations stay labelled. Illustrative, not a prediction."

Extras, only after the main answer finishes: "Should Messi press earlier?" (no such player) and "Should Kiprono make a run or hold?" (only the right-back's press can be simulated). Both are instant.