Second Look — demo video recording sheet

Target **1:50** (hard limit under 2:00). Upload to YouTube as **Public**. Deadline: 27 Oct, 11:59 pm Pacific (2:59 am Eastern).

Two takes, one edit

• **Take A — hosted site** https://second-look-dusky.vercel.app (offline mode, every click instant): The press, Three choices, Try to flip it, On air.

• **Take B — local app** http://localhost:5173 with the backend in on-device mode (Foundry Local, qwen3-8b): The argument with the live agents, then Fan view. One answer takes 1–3 minutes: speed it up in the edit and label it **"≈2 min later — sped up"**.

• Record with Clipchamp (Record & create → Screen) or Xbox Game Bar (Win+Alt+R). Export raw takes as 1080p MP4 to `C:\Users\benon\Videos\SecondLook` as takeA.mp4 and takeB.mp4.

Before recording

• Browser about 1600×1000, zoom 100%, bookmarks bar hidden, Do Not Disturb on, other tabs closed.

• Take B: `foundry server start`, `foundry model load qwen3-8b`, restart the backend so it runs the latest code; header must read *Foundry Local (on-device) · qwen3-8b · Agent Framework*. Ask one warm-up question, then refresh.

• 10-second mic test. Speak slowly; retakes per row are fine.

Don't

• Say the hosted site uses a live model, or call the on-device model "cloud Foundry".

• Claim personalization, real-time analysis, multiple languages, real players or real-match accuracy.

• Hide the on-device wait or the "updating" state; no club logos, real footage or copyrighted music; don't end on the JSON download.

Upload

• YouTube → **Public**. Title: *Second Look — the football argument, with its assumptions on the pitch*.

• Description: first paragraph of docs/SUBMISSION.md, then https://github.com/Ben369369/Second-Look and https://second-look-dusky.vercel.app. Check it plays signed out.

The script

Time	Take	On screen	Say (suggested)
0:00–0:15	A	**The press** opens on the question, the verdict "No. Holding was safer.", 30 vs 71, and the pitch frozen with the orange 8.5 m gap.	"Should the right-back have pressed earlier? Second Look replays the next eight seconds two hundred times for each choice. If he holds, Harbor reach a dangerous position in 30 of 200. If he presses earlier, 71. Lower is safer: holding wins."
0:15–0:38	A	**Three choices.** Drag **Cover delay** 0.6 s → 1.2 s; wait for "updating" to clear; the headline becomes "Holding position is the clear call".	"With a teammate covering the inside pass it's too close to call, under a rule we fixed before any run. Make that cover arrive 0.6 seconds later and holding becomes the clear call. The answer depends on that assumption, so the app shows it."
0:38–0:45	A	Under **Try to flip it**, click **Run stress test**; show the sentence.	"A challenger searches a grid of assumptions for settings that flip the answer."
0:45–1:15	B	**The argument**, Analyst view. Ask the pressing question; show "planning…". Cut, label "≈2 min later — sped up". Show **AI agents · on-device · qwen3-8b**; expand **How this answer was made**; scroll to the validator line with a rejection.	"Here the same question goes to four Microsoft Agent Framework agents — a planner, an evidence analyst, a challenger and a producer — running on this laptop with Foundry Local. The agents choose the evidence and explain it, but code writes every number. When the producer wrote a claim the evidence can't support, the validator rejected it."
1:15–1:27	B	Switch to **Fan view** (instant).	"Same facts for a fan: plain words, the caveat attached, nothing re-run."
1:27–1:48	A or B	**On air.** Click **Play recap**; let a card appear on the pitch; end on the SIMULATION card.	"For a broadcast desk, the answer becomes timed cards on the replay. Every card cites a fact, and simulations stay labelled. Illustrative, not a prediction: an argument with its assumptions out in the open."

If the on-device run has no rejected claim, point at the validator line ("N approved, M rejected") and say "a validator checks every claim" instead.